





## INDEX

Sr No.	Name of the practical	Date	Signature
1	Install, configure and run Hadoop and HDFS and explore HDFS.		
2	Implement word count / frequency programs using MapReduce		
3	Implement an application that stores big data in Hbase/MongoDB and manipulate it using R/Python		
4	Write a Pig Script for solving counting problems.		
5	Configure the Hive and implement the application in Hive.		
6	Implement the following		
6A	Implement Decision tree classification techniques		
6B	Implement SVM classification techniques		
7	Solve the following:		
7A	REGRESSION MODEL Import a data from web storage. Name the dataset and now do Logistic Regression to find out relation between variables that are affecting the admission of a student in an institute based on his or her GRE score, GPA obtained and rank of the student. Also check the model is fit or not. require (foreign), require(MASS).		
7B	MULTIPLE REGRESSION MODEL Apply multiple regressions, if data have a		

	continuous independent variable. Apply on above dataset.		
8	Solve the Following:		
8A	<b>CLASSIFICATION MODEL</b> a. Install relevant package for classification. b. Choose classifier for classification problem. c. Evaluate the performance of classifier.		
8B	<b>CLUSTERING MODEL</b> a. Clustering algorithms for unsupervised classification. b. Plot the cluster data using R visualizations.		

# PRACTICAL 1

## Hadoop Installation

### Aim:

Install, configure and run Hadoop and HDFS and explore HDFS on Windows

### Code:

#### Steps to Install Hadoop

1. Install Java JDK 1.8
2. Download Hadoop and extract and place under C drive
3. Set Path in Environment Variables
4. Config files under Hadoop directory
5. Create folder datanode and namenode under data directory
6. Edit HDFS and YARN files
7. Set Java Home environment in Hadoop environment
8. Setup Complete. Test by executing start-all.cmd

There are two ways to install Hadoop, i.e.

9. Single node

10. Multi node

Here, we use multi node cluster.

#### 1. Install Java

11. – Java JDK Link to download

<https://www.oracle.com/java/technologies/javase-jdk8-downloads.html>

12. – extract and install Java in C:\Java

13. – open cmd and type -> javac -version

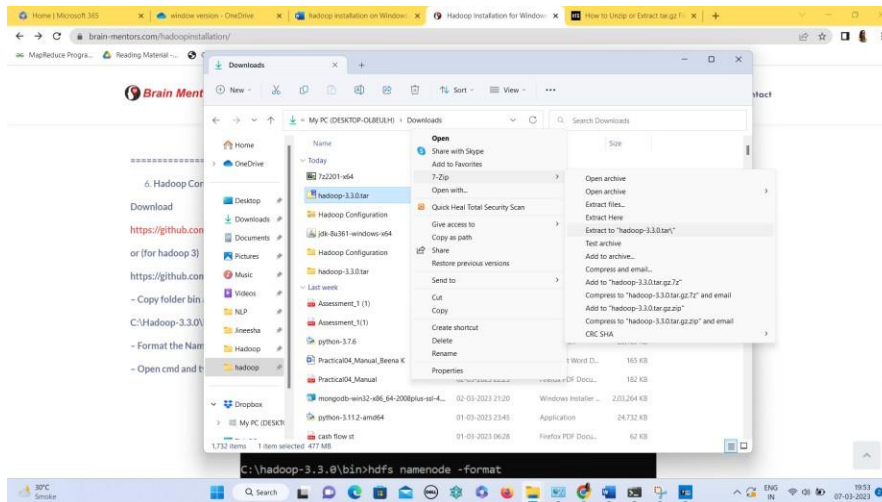
```
C:\Users>cd Beena

C:\Users\Beena>java -version
java version "1.8.0_361"
Java(TM) SE Runtime Environment (build 1.8.0_361-b09)
Java HotSpot(TM) 64-Bit Server VM (build 25.361-b09, mixed mode)
```

#### 2. Download Hadoop

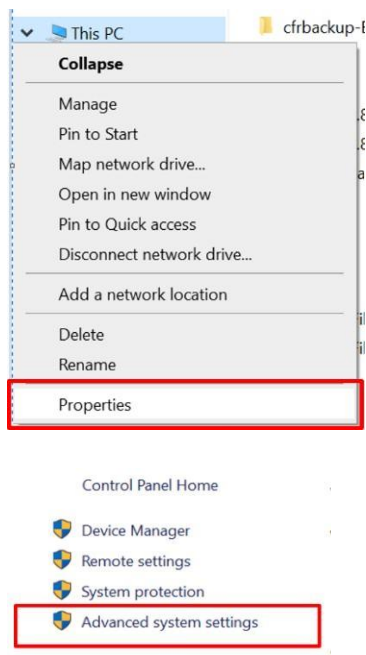
<https://www.apache.org/dyn/closer.cgi/hadoop/common/hadoop-3.3.0/hadoop-3.3.0.tar.gz>

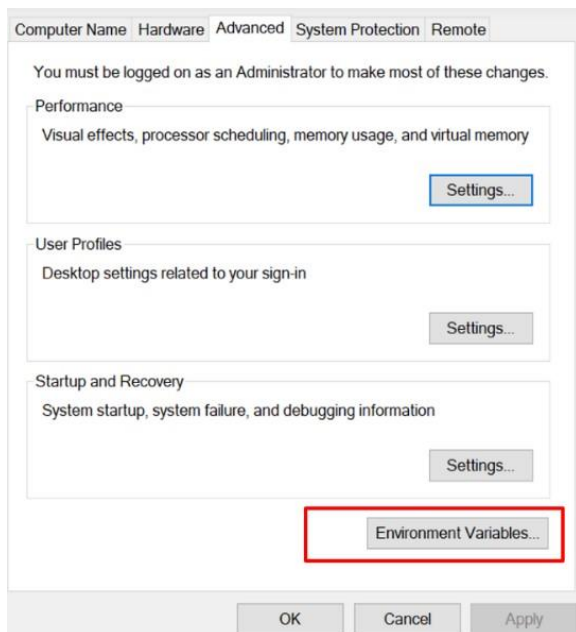
- right click .rar.gz file -> show more options -> 7-zip->and extract to C:\Hadoop-3.3.0\



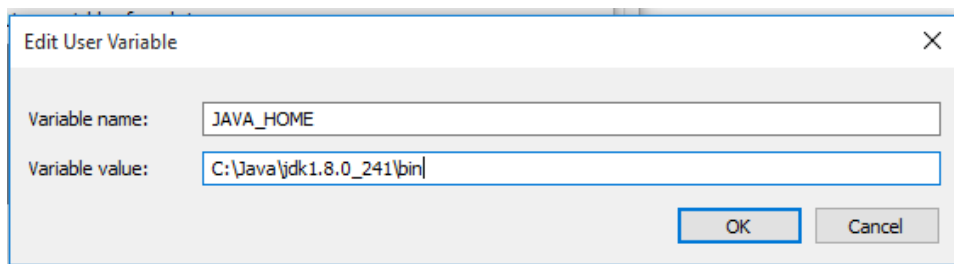
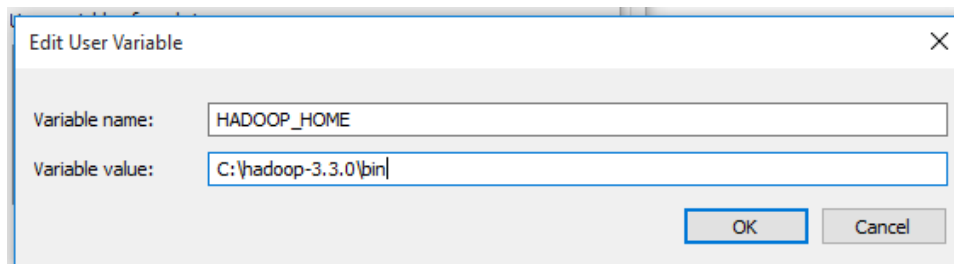
**3. Set the path JAVA\_HOME Environment variable**

**4. Set the path HADOOP\_HOME Environment variable**

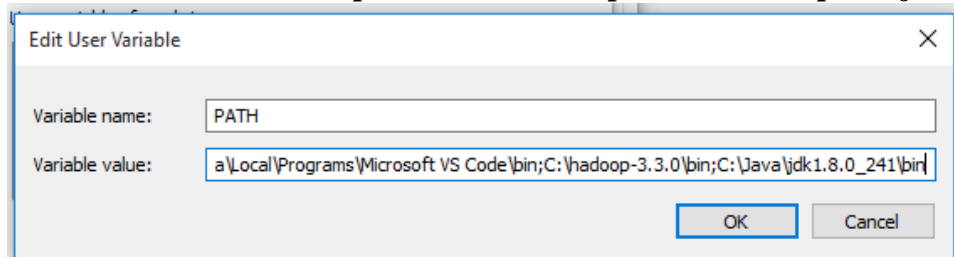




Click on **New** to both user variables and system variables.



Click on user variable -> path -> edit-> add path for Hadoop and java upto 'bin'



Click Ok, Ok, Ok.

## 5. Configurations

Edit file C:/Hadoop-3.3.0/etc/hadoop/core-site.xml,

paste the xml code in folder and save

```
=====

<configuration>
<property>
  <name>fs.defaultFS</name>
  <value>hdfs://localhost:9000</value>
</property>
</configuration>

=====
```

**Rename “mapred-site.xml.template” to “mapred-site.xml” and edit this file C:/Hadoop-3.3.0/etc/hadoop/mapred-site.xml, paste xml code and save this file.**

```
=====

<configuration>
<property>
  <name>mapreduce.framework.name</name>
  <value>yarn</value>
</property>
</configuration>

=====
```

**Create folder “data” under “C:\Hadoop-3.3.0”**

**Create folder “datanode” under “C:\Hadoop-3.3.0\data”**

**Create folder “namenode” under “C:\Hadoop-3.3.0\data”**

**Edit file C:\Hadoop-3.3.0/etc/hadoop/hdfs-site.xml,**

**paste xml code and save this file.**

```
=====

<configuration>
<property>
  <name>dfs.replication</name>
  <value>1</value>

```



```
</property>
```

```
<property>
```

```
  <name>dfs.namenode.name.dir</name>
```

```
  <value>/hadoop-3.3.0/data/namenode</value>
```

```
</property>
```

```
<property>
```

```
  <name>dfs.datanode.data.dir</name>
```

```
  <value>/hadoop-3.3.0/data/datanode</value>
```

```
</property>
```

```
</configuration>
```

=====

**Edit file C:/Hadoop-3.3.0/etc/hadoop/yarn-site.xml,**

**paste xml code and save this file.**

```
<configuration>
```

```
  <property>
```

```
    <name>yarn.nodemanager.aux-services</name>
```

```
    <value>mapreduce_shuffle</value>
```

```
  </property>
```

```
  <property>
```

```
    <name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</name>
```

```
    <value>org.apache.hadoop.mapred.ShuffleHandler</value>
```

```
  </property>
```

```
  <property>
```

```
    <name>yarn.resourcemanager.address</name>
```

```
    <value>127.0.0.1:8032</value>
```

```
  </property>
```

```
  <property>
```

```
    <name>yarn.resourcemanager.scheduler.address</name>
```

```

        <value>127.0.0.1:8030</value>
    </property>
    <property>
        <name>yarn.resourcemanager.resource-tracker.address</name>
        <value>127.0.0.1:8031</value>
    </property>
</configuration>

```

## 6. Edit file C:/Hadoop-3.3.0/etc/hadoop/hadoop-env.cmd

Find “JAVA\_HOME=%JAVA\_HOME%” and replace it as  
 set JAVA\_HOME="C:\Java\jdk1.8.0\_361"

## 7. Download “redistributable” package

### Download and run VC\_redist.x64.exe

This is a “redistributable” package of the Visual C runtime code for 64-bit applications, from Microsoft. It contains certain shared code that every application written with Visual C expects to have available on the Windows computer it runs on.

## 8. Hadoop Configurations

Download **bin** folder from

<https://github.com/s911415/apache-hadoop-3.1.0-winutils>

– Copy the **bin** folder to c:\hadoop-3.3.0. Replace the existing **bin** folder.

9. copy "hadoop-yarn-server-timelineservice-3.0.3.jar" from ~\hadoop-3.0.3\share\hadoop\yarn\timelineservice to ~\hadoop-3.0.3\share\hadoop\yarn folder.

## 10. Format the NameNode

– Open cmd ‘Run as Administrator’ and type command “hdfs namenode –format”

```
C:\> Administrator: Command Prompt

Microsoft Windows [Version 10.0.22621.1265]
(c) Microsoft Corporation. All rights reserved.

C:\Windows\System32>cd\hadoop-3.3.0\bin

C:\hadoop-3.3.0\bin>hdfs namenode -format
```

```
2023-03-07 21:31:34,685 INFO namenode.FSImageFormatProtobuf: Saving image file C:\hadoop-3.3.0\data\namenode\current\fsi
image.ckpt_00000000000000000000 using no compression
2023-03-07 21:31:34,844 INFO namenode.FSImageFormatProtobuf: Image file C:\hadoop-3.3.0\data\namenode\current\fsimage.ck
pt_00000000000000000000 of size 400 bytes saved in 0 seconds .
2023-03-07 21:31:34,860 INFO namenode.NNStorageRetentionManager: Going to retain 1 images with txid >= 0
2023-03-07 21:31:34,869 INFO namenode.FSImage: FSImageSaver clean checkpoint: txid=0 when meet shutdown.
2023-03-07 21:31:34,870 INFO namenode.NameNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down NameNode at DESKTOP-OL8EULH/192.168.1.19
*****/
```

## 11. Testing

- Open cmd ‘Run as Administrator’ and change directory to C:\Hadoop-3.3.0\sbin
- type start-all.cmd

OR

- type start-dfs.cmd
- type start-yarn.cmd

```
C:\hadoop-3.3.0\sbin>start-all.cmd
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
The filename, directory name, or volume label syntax is incorrect.
The filename, directory name, or volume label syntax is incorrect.
starting yarn daemons
The filename, directory name, or volume label syntax is incorrect.
```

- You will get 4 more running threads for Datanode, namenode, resource manager and node manager

```
Apache Hadoop Distribution - hadoop namenode
2023-03-07 20:33:00,395 INFO ipc.Server: Starting Socket Reader #1 for port 9000
2023-03-07 20:33:00,547 INFO namenode.FSNamesystem: Registered FSNamesystemState, ReplicatedBlocksState and ECBlockGroup
sState MBeans.
2023-03-07 20:33:00,549 INFO common.Util: Assuming 'file' scheme for path /hadoop-3.3.0/data/namenode in configuration.
2023-03-07 20:33:00,554 INFO namenode.LeaseManager: Number of blocks under construction: 0
2023-03-07 20:33:00,563 INFO blockmanagement.DatanodeAdminDefaultMonitor: Initialized the Default Decommission and Mainte
enance monitor
2023-03-07 20:33:00,566 INFO blockmanagement.BlockManager: initializing replication queues
2023-03-07 20:33:00,567 INFO hdfs.StateChange: STATE* Leaving safe mode after 0 secs
2023-03-07 20:33:00,567 INFO hdfs.StateChange: STATE* Network topology has 0 racks and 0 datanodes
2023-03-07 20:33:00,569 INFO hdfs.StateChange: STATE* UnderReplicatedBlocks has 0 blocks
2023-03-07 20:33:00,574 INFO blockmanagement.BlockManager: Total number of blocks = 0
2023-03-07 20:33:00,575 INFO blockmanagement.BlockManager: Number of invalid blocks = 0
2023-03-07 20:33:00,575 INFO blockmanagement.BlockManager: Number of under-replicated blocks = 0
2023-03-07 20:33:00,576 INFO blockmanagement.BlockManager: Number of over-replicated blocks = 0
2023-03-07 20:33:00,576 INFO blockmanagement.BlockManager: Number of blocks being written = 0
2023-03-07 20:33:00,576 INFO hdfs.StateChange: STATE* Replication Queue initialization scan for invalid, over- and under
-replicated blocks completed in 9 msec
2023-03-07 20:33:00,607 INFO ipc.Server: IPC Server Responder: starting
2023-03-07 20:33:00,607 INFO ipc.Server: IPC Server listener on 9000: starting
2023-03-07 20:33:00,611 INFO namenode.NameNode: NameNode RPC up at: localhost/127.0.0.1:9000
2023-03-07 20:33:00,614 INFO namenode.FSNamesystem: Starting services required for active state
2023-03-07 20:33:00,614 INFO namenode.FSDirectory: Initializing quota with 4 thread(s)
2023-03-07 20:33:00,622 INFO namenode.FSDirectory: Quota initialization completed in 7 milliseconds
name space=1
storage space=0
storage types=RAM_DISK=0, SSD=0, DISK=0, ARCHIVE=0, PROVIDED=0
2023-03-07 20:33:00,626 INFO blockmanagement.CacheReplicationMonitor: Starting CacheReplicationMonitor with interval 300
00 milliseconds
```

**Output:**

12. Type JPS command to start-all.cmd command prompt, you will get following output.

```
C:\hadoop-3.3.0\sbin>jps
5632 Jps
7572 DataNode
3752 ResourceManager
7992 NameNode
8028 NodeManager
```

13. Run <http://localhost:9870/> from any browser

The screenshot shows a web browser window with the address bar displaying `localhost:9870/dfshealth.html#tab-overview`. The page has a green header with navigation tabs: **Hadoop**, **Overview**, **Datanodes**, **Datanode Volume Failures**, **Snapshot**, **Startup Progress**, and **Utilities**. The main content area is titled **Overview 'localhost:9000' (✓active)**. Below the title is a table with the following information:

Started:	Wed Mar 15 12:10:54 +0530 2023
Version:	3.3.0, raa96f1871bfd858f9bac59cf2a81ec470da649af
Compiled:	Tue Jul 07 00:14:00 +0530 2020 by brahma from branch-3.3.0
Cluster ID:	CID-1986aba8-0ed3-43a2-9db7-42944ec518b2
Block Pool ID:	BP-1049743432-192.168.56.1-1678862097216

## Browse Directory

Go!

Show25 ▾ entries

Search:

☐

Permission

Owner

Group

Size

Last Modified

Replication

Block Size

Name

No data available in table

Showing 0 to 0 of 0 entries

Previous

Next

## PRACTICAL 2

### MapReduce Implementation

#### Aim:

Implement word count / frequency programs using MapReduce.

#### Steps:

```
C:\hadoop-3.3.0\sbin>start-dfs.cmd
```

```
C:\hadoop-3.3.0\sbin>start-yarn.cmd
```

Open a command prompt as administrator and run the following command to create an input and output folder on the Hadoop file system, to which we will be moving the sample.txt file for our analysis.

```
C:\hadoop-3.3.0\bin>cd\
```

```
C:\>hadoop dfsadmin -safemode leave
```

DEPRECATED: Use of this script to execute hdfs command is deprecated.

Instead use the hdfs command for it.

Safe mode is OFF

```
C:\>hadoop fs -mkdir /input_dir
```





Check it by giving the following URL at browser

<http://localhost:9870>


Utilities -> browse the file system

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

### Browse Directory

/ Go!    

Show 25 entries Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	<input type="checkbox"/>
<input type="checkbox"/>	drwxr-xr-x	vinod	supergroup	0 B	Mar 30 19:08	0	0 B	input_dir	

Showing 1 to 1 of 1 entries

Previous 1 Next

Copy the input text file named input\_file.txt in the input directory (input\_dir) of HDFS.

Make a file in c:\input\_file.txt and write following content in it.

Hadoop Window version is easy compared to Ubuntu version

Now apply the following command at c:\>

```
C:\> hadoop fs -put C:/input_file.txt /input_dir
```

## Browse Directory

/input\_dir

Go!

Show

25

entries

Search:

<div><input type="checkbox"/></div>	<div><div><div></div><div></div></div>Permission</div>	<div><div><div></div><div></div></div>Owner</div>	<div><div><div></div><div></div></div>Group</div>	<div><div><div></div><div></div></div>Size</div>	<div><div><div></div><div></div></div>Last Modified</div>	<div><div><div></div><div></div></div>Replication</div>	<div><div><div></div><div></div></div>Block Size</div>	<div><div><div></div><div></div></div>Name</div>	<div><div><div></div><div></div></div></div>
<div><input type="checkbox"/></div>	<div><div><div></div><div></div></div>-rw-r--r--</div>	<div><div><div></div><div></div></div>vinod</div>	<div><div><div></div><div></div></div>supergroup</div>	<div><div><div></div><div></div></div>57 B</div>	<div><div><div></div><div></div></div>Mar 30 19:11</div>	<div><div><div></div><div></div></div>1</div>	<div><div><div></div><div></div></div>128 MB</div>	<div><div><div></div><div></div></div>input_file.txt</div>	<div><div><div></div><div></div></div></div>

Showing 1 to 1 of 1 entries

Previous

1

Next

### Verify input\_file.txt available in HDFS input directory (input\_dir).

C:\>Hadoop fs -ls /input\_dir/

```
C:\>hadoop fs -put C:/input_file.txt /input_dir

C:\>hadoop fs -ls /input_dir/
Found 1 items
-rw-r--r--   1 vinod supergroup          57 2023-03-30 19:11 /input_dir/input_file.txt

C:\>
```

### Verify content of the copied file

C:\>hadoop dfs -cat /input\_dir/input\_file.txt

You can see the file content displayed on the CMD.

```
C:\>hadoop dfs -cat /input_dir/input_file.txt
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.
Hadoop Window version is easy compared to Ubuntu version.

C:\>
```

Run MapReduceClient.jar and also provide input and out directories.

C:\>hadoop jar C:/hadoop-3.3.0/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.3.0.jar wordcount /input\_dir /output\_dir

```
Reduce input groups=8
Reduce shuffle bytes=103
Reduce input records=8
Reduce output records=8
Spilled Records=16
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=70
CPU time spent (ms)=219
Physical memory (bytes) snapshot=517128192
Virtual memory (bytes) snapshot=792633344
Total committed heap usage (bytes)=392691712
Peak Map Physical memory (bytes)=314761216
Peak Map Virtual memory (bytes)=465485824
Peak Reduce Physical memory (bytes)=202366976
Peak Reduce Virtual memory (bytes)=327180288

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=56
File Output Format Counters
  Bytes Written=65

C:\Windows\System32>
```

In case, there is some error in executing then copy the file MapReduceClient.jar in C:\ and run the program with the jar file using existing MapReduceClient.jar file as:

C:\> hadoop jar C:/MapReduceClient.jar wordcount /input\_dir /output\_dir

Now, check the output\_dir on browser as follows:

**Hadoop** Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

### Browse Directory

Show  entries Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	drwxr-xr-x	vinod	supergroup	0 B	Mar 30 19:29	<u>0</u>	0 B	<a href="#">input_dir</a>	
<input type="checkbox"/>	drwxr-xr-x	vinod	supergroup	0 B	Mar 30 19:30	<u>0</u>	0 B	<a href="#">output_dir</a>	

Click on output\_dir ☐ part-r-00000 ☐ Head the file (first 32 K) and check the file content as the output.



File information - part-r-00000

Download

Head the file (first 32K)

Tail the file (last 32K)

Block information — Block 0

Block ID: 1073741832  
Block Pool ID: BP-537931513-192.168.1.19-1680861805234  
Generation Stamp: 1008  
Size: 65  
Availability:

- DESKTOP-OL8EULH

File contents

Hadoop 1  
Ubuntu 1  
Window 1  
compared 1  
easy 1  
is 1  
to 1  
version 2

Search:

Block Size

Name

MB

\_SUCCESS

MB

part-r-00000

Previous

1

Next

Alternatively, you may type the following command on CMD window as:

```
C:\> hadoop dfs -cat /output_dir/*
```

You can get the following output

```
C:\Windows\System32>hadoop dfs -cat /output_dir/*
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.
Hadoop 1
Ubuntu 1
Window 1
compared 1
easy 1
is 1
to 1
version 2
C:\Windows\System32>
```

## PRACTICAL 3

### Aim:

Implement an application that stores big data in Hbase / MongoDB and manipulate it using R / Python

### Requirements

- a. PyMongo
- b. Mongo Database

### Step A: Install Mongo database

Step 1) Go to (<https://www.mongodb.com/download-center/community>) and Download MongoDB Community Server. We will install the 64-bit version for Windows.

Select the server you would like to run:



MongoDB Community Server  
FEATURE RICH. DEVELOPER READY.

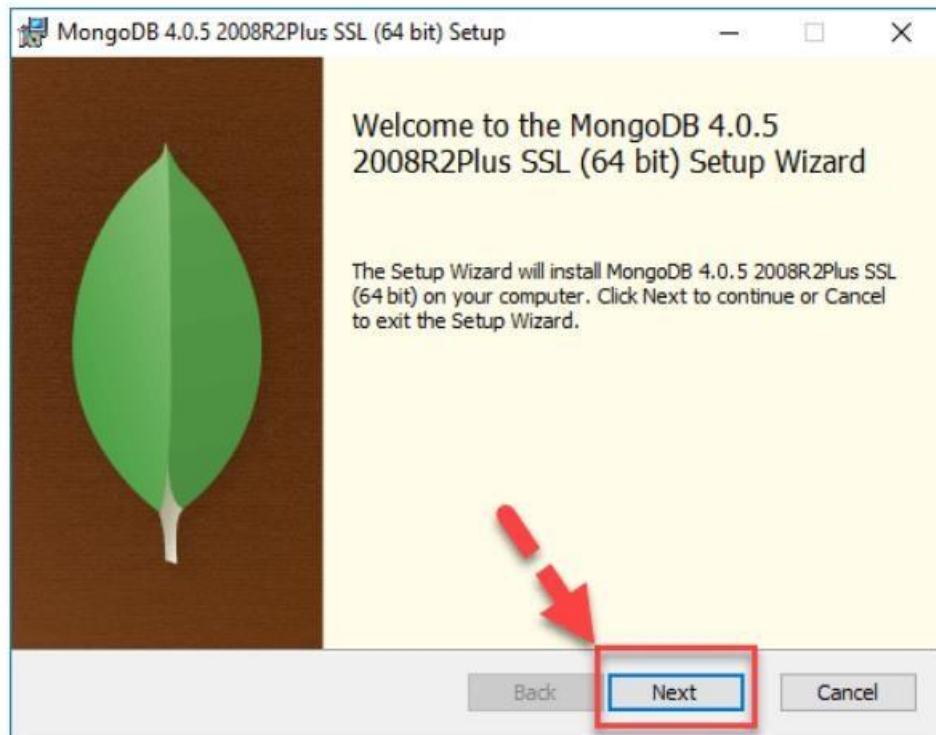
Version: 4.0.5 (current release) OS: Windows 64-bit x64

Package: MSI

**Download**

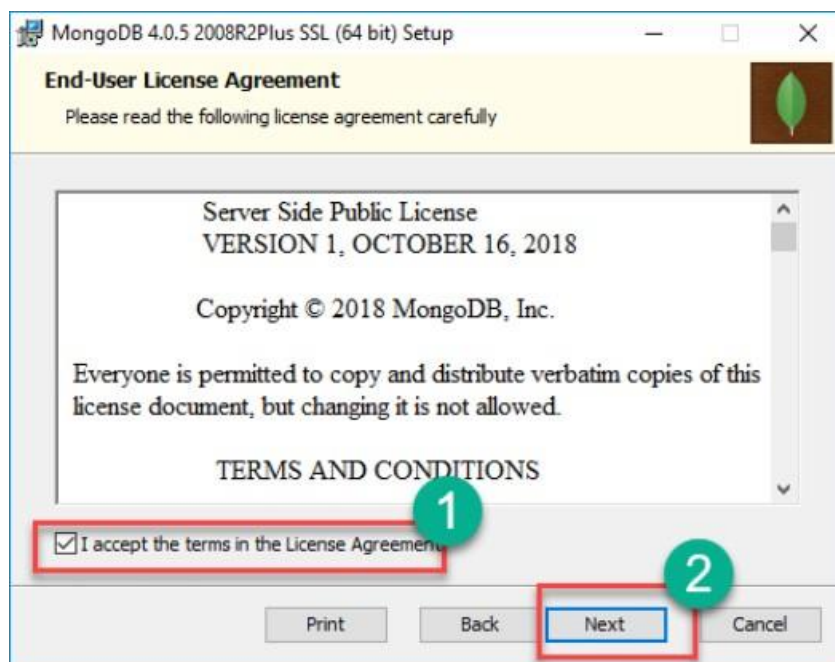
[https://fastdl.mongodb.org/win32/mongodb-win32-x86\\_64-2008plus-ssl-4.0.5-signed.msi](https://fastdl.mongodb.org/win32/mongodb-win32-x86_64-2008plus-ssl-4.0.5-signed.msi)

Step 2) Once download is complete open the msi file. Click Next in the start up screen




Step 3)

1. Accept the End-User License Agreement
2. Click Next

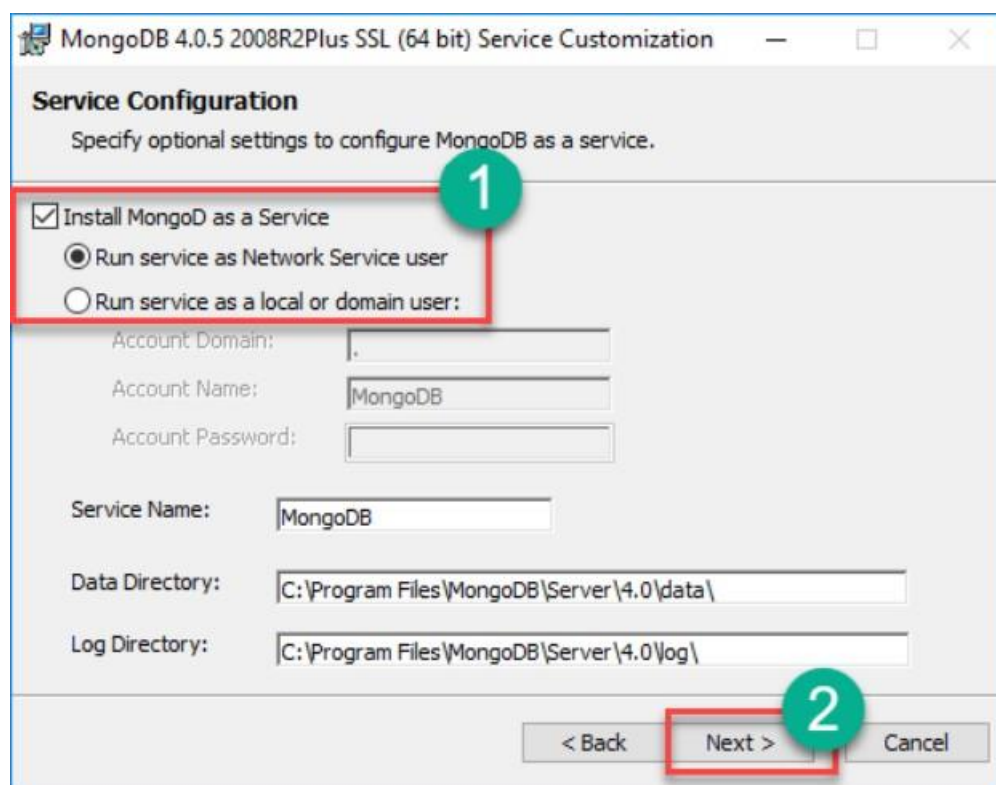


Step 4) Click on the "complete" button to install all of the components. The custom option can be used to install selective components or if you want to change the location of the installation.

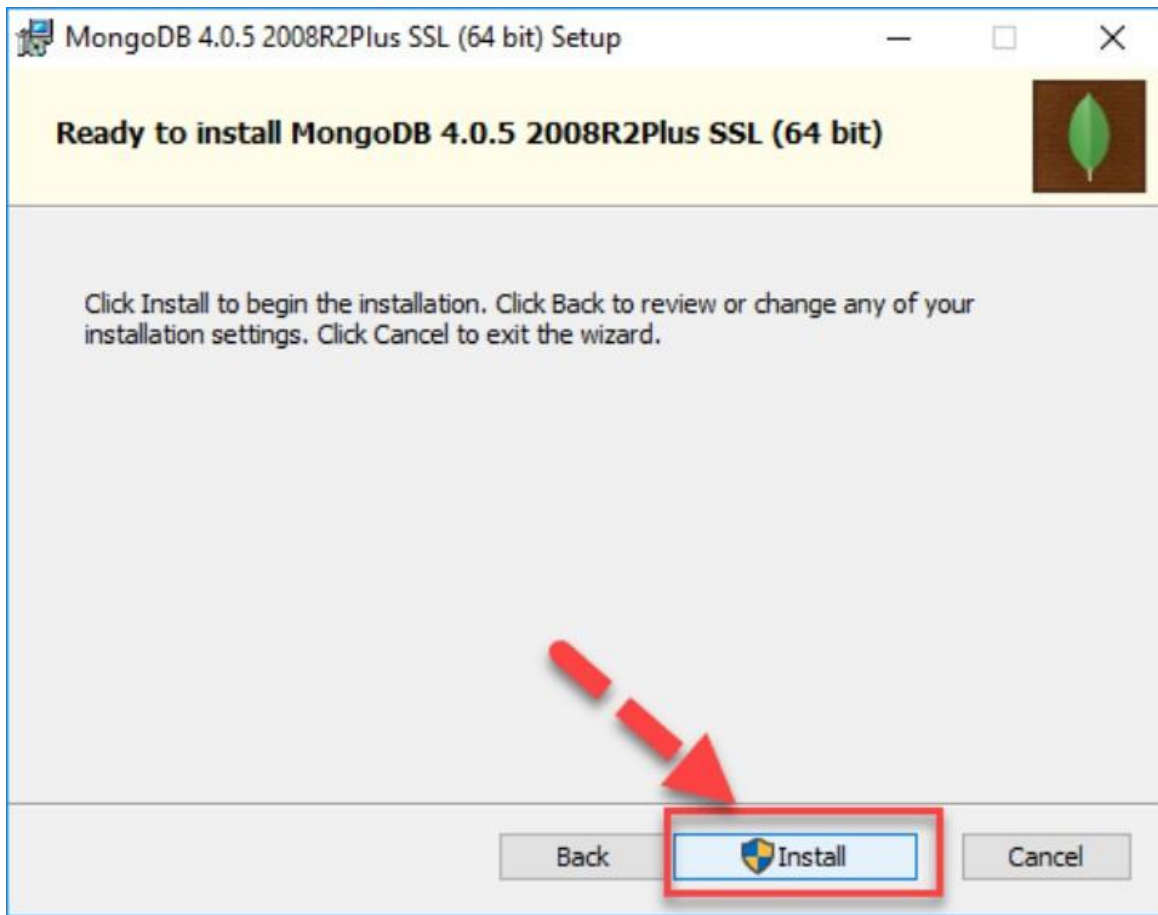


Step 5)


1. Select “Run service as Network Service user”. make a note of the data directory, we’ll need this later.
2. Click Next



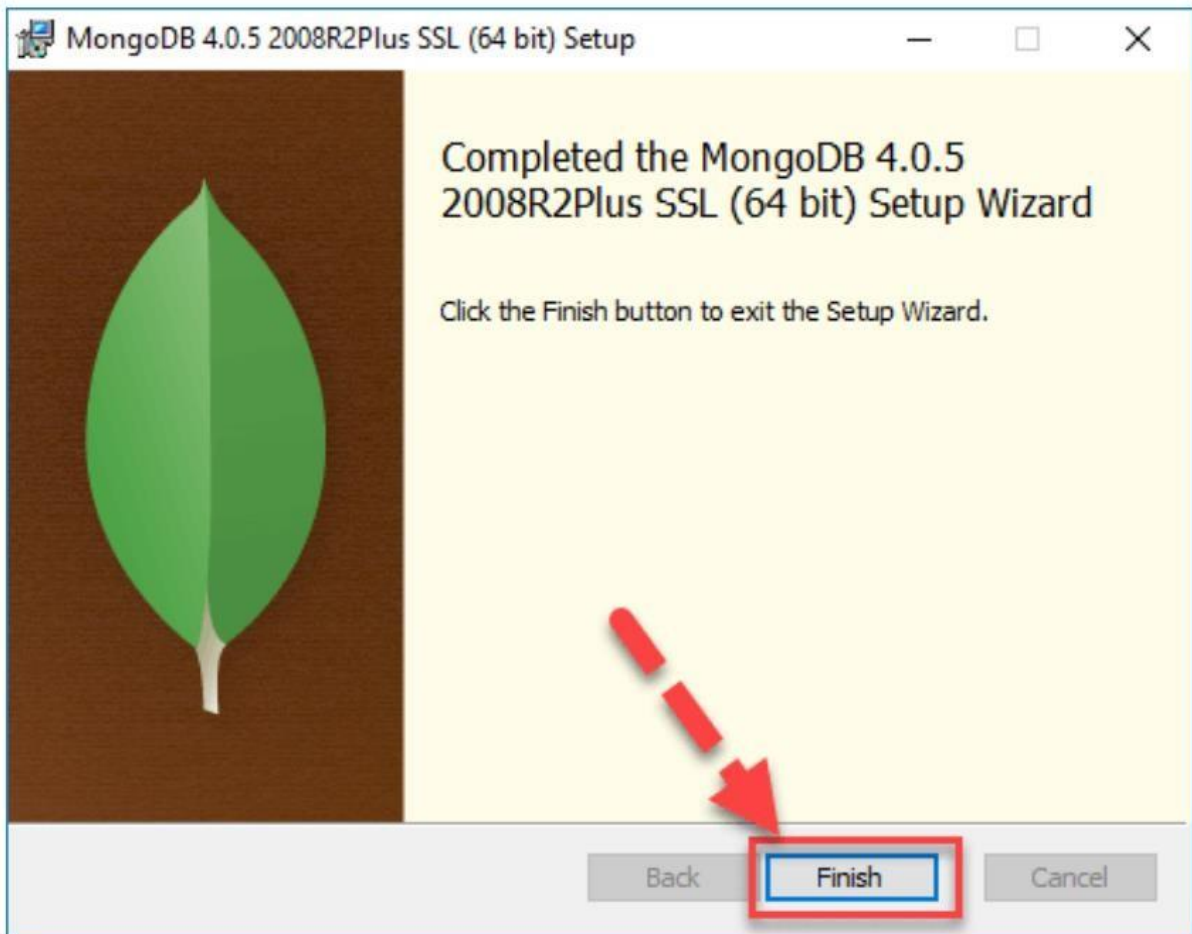
Step 6) Click on the Install button to start the installation.



Step 7) Installation begins. Click Next once completed.



Step 8) Click on the Finish button to complete the installation.



### Test MongoDB

**Step 1)** Go to " C:\Program Files\MongoDB\Server\4.0\bin" and double click on **mongo.exe**. Alternatively, you can also click on the MongoDB desktop icon.

- **Create the directory where MongoDB will store its files.**  
Open command prompt window and apply following commands

```
C:\users\admin> cd\  
C:\>md data\db
```

### Step 2) Execute mongod

Open another command prompt window.

```
C:\> cd C:\Program Files\MongoDB\Server\4.0\bin  
C:\Program Files\MongoDB\Server\4.0\bin> mongod
```

*In case if it gives an error then run the following command:*

```
C:\Program Files\MongoDB\Server\4.0\bin> mongod -repair
```

```
{ v: 2, key: { lastUse: 1 }, name: "lsidTTLIndex", ns: "config.system.sessions", expireAfterSeconds: 1800 }
2023-03-03T09:46:21.011+0530 I INDEX [LogicalSessionCacheRefresh] building index using bulk method; build may temporarily use up to 500 megabytes of RAM
2023-03-03T09:46:21.044+0530 I INDEX [LogicalSessionCacheRefresh] build index done. scanned 0 total records. 0 secs
2023-03-03T09:46:21.045+0530 I COMMAND [LogicalSessionCacheRefresh] command config.$cmd command: createIndexes { createIndexes: "system.sessions", indexes: [ { key: { lastUse: 1 }, name: "lsidTTLIndex", expireAfterSeconds: 1800 } ], $db: "config" } numYields:0 reslen:114 locks:{ Global: { acquireCount: { r: 2, w: 2 } }, Database: { acquireCount: { w: 2, W: 1 } }, Collection: { acquireCount: { w: 2 } } } protocol:op_msg 254ms
```

### Step 3) Connect to MongoDB using the Mongo shell

Let the MongoDB daemon to run.

Open another command prompt window and run the following commands:

C:\users\admin> cd C:\Program Files\MongoDB\Server\4.0\bin

C:\Program Files\MongoDB\Server\4.0\bin>mongo

```
The monitoring data will be available on a MongoDB website with a unique URL accessible to you
and anyone you share the URL with. MongoDB may use this information to make product
improvements and to suggest MongoDB products and deployment options to you.

To enable free monitoring, run the following command: db.enableFreeMonitoring()
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()
---

> show dbs
admin      0.000GB
config     0.000GB
local      0.000GB
mybigdata  0.000GB
> use mybigdata
```

### Step 4) Install PyMongo

Open another command prompt window and run the following commands:

Check the python version on your desktop / laptop and copy that path from window explorer

C:\users\admin>cd C:\Program Files\Python311\Scripts

C:\Program Files\Python38\Scripts > python -m pip install pymongo

```
C:\Program Files\Python38\Scripts>python -m pip install pymongo
Defaulting to user installation because normal site-packages is not writeable
Collecting pymongo
  Downloading pymongo-4.3.3-cp38-cp38-win_amd64.whl (382 kB)
    |#####| 382 kB 3.3 MB/s
Collecting dnspython<3.0.0,>=1.16.0
  Downloading dnspython-2.3.0-py3-none-any.whl (283 kB)
    |#####| 283 kB ...
Installing collected packages: dnspython, pymongo
Successfully installed dnspython-2.3.0 pymongo-4.3.3
WARNING: You are using pip version 20.2.1; however, version 23.0.1 is available.
You should consider upgrading via the 'C:\Program Files\Python38\python.exe -m pip install --upgrade pip' command.
```

Note: # -m option is for <module-name>

Now you have downloaded and installed a mongoDB driver.

### Step 5) Test PyMongo

Run the following command from python command prompt

```
import pymongo
```

Now, either create a file in Python IDLE or run all commands one by one in sequence on Python cell

#### Program 1: Creating a Database: create\_dp.py

```
import pymongo
```



```
myclient = pymongo.MongoClient("mongodb://localhost:27017/")
mydb = myclient["mybigdata"]
print(myclient.list_database_names())
['admin', 'config', 'local']
```

### **Program 2: Creating a Collection: create\_collection.py**

```
import pymongo
myclient = pymongo.MongoClient("mongodb://localhost:27017/")
mydb = myclient["mybigdata"]
mycol=mydb["student"]
print(mydb.list_collection_names())
[]
```

### **Program 3: Insert into Collection: insert\_into\_collection.py**

```
import pymongo
myclient = pymongo.MongoClient("mongodb://localhost:27017/")
mydb = myclient["mybigdata"]
mycol=mydb["student"]
mydict={"name":"Beena", "address":"Mumbai"}
x=mycol.insert_one(mydict) # insert_one(containing the name(s) and value(s) of each field
```

### **Program 4: Insert Multiple data into Collection: insert\_many.py**

```
import pymongo
myclient = pymongo.MongoClient("mongodb://localhost:27017/")
mydb = myclient["mybigdata"]
mycol=mydb["student"]
mylist=[{"name":"Khyati", "address":"Mumbai"}, {"name":"Kruti", "address":"Mumbai"},
{"name":"Nidhi", "address":"Pune"}, {"name":"Komal", "address":"Pune"},]
x=mycol.insert_many(mylist)
```

### **Step 6) Test in Mongodb to check database and data inserted in collection**

a. If you want to check your database list, use the command show dbs in mongo command prompt

> show dbs

```
admin      0.000GB
config     0.000GB
local      0.000GB
mybigdata  0.000GB
```

b. If you want to use a database with name mybigdata, then use database statement would be as follow:

> use mybigdata

```
switched to db mybigdata
```



c. If you want to check collection in mongodb use the command show collections

> show collections

```
student
```

d. If you want to display the first row from collection: db.collection\_name.findOne()

> db.student.findOne()

```
> db.student.findOne()
{
  "_id" : ObjectId("640178face663db608cef72f"),
  "name" : "Beena",
  "address" : "Mumbai"
}
```

e. If you want to display all the data from collection: db.collection\_name.find()

> db.student.find()

```
> db.student.find()
{ "_id" : ObjectId("640178face663db608cef72f"), "name" : "Beena", "address" : "Mumbai" }
{ "_id" : ObjectId("640179336ce317082c266dc1"), "name" : "Khyati", "address" : "Mumbai" }
{ "_id" : ObjectId("640179336ce317082c266dc2"), "name" : "Kruti", "address" : "Mumbai" }
{ "_id" : ObjectId("640179336ce317082c266dc3"), "name" : "Nidhi", "address" : "Pune" }
{ "_id" : ObjectId("640179336ce317082c266dc4"), "name" : "Komal", "address" : "Pune" }
```

f. count number of rows in a collection

> db.student.count()

```
5
```

**Site for R packages documentation:**

[https://cran.r-project.org/web/packages/available\\_packages\\_by\\_name.html](https://cran.r-project.org/web/packages/available_packages_by_name.html)

## PRACTICAL 4

**Aim :** Write a Pig Script for solving counting problems.

**Steps :**

```
cat> /home/cloudera/input.csv
```

```
cat /home/cloudera/input.csv
```

```
pig -x local
```

```
lines = load '/home/cloudera/input.csv' as (line:chararray);
```

```
words = foreach lines GENERATE FLATTEN(TOKENIZE(line)) as word;
```

```
grouped = GROUP words by word;
```

```
wordcount = foreach grouped GENERATE group, COUNT(words);
```

```
dump wordcount;
```

## PRACTICAL 5

**Aim:** Install Hive and use Hive Create and store structured databases.

**Steps:**

```
cat > /home/cloudera/employee.txt
```

```
1~Sachine~Pune~Product Engineering~100000~Big Data
```

```
2~Gaurav~Banglore~Sales~90000~CRM
```

```
3~Manish~Chennai~Recruiter~125000~HR
```

```
4~Bhushan~Hyderabad~Developer~50000~BFSI
```

```
cat /home/cloudera/employee.txt
```

```
sudo -u hdfs hadoop fs -put /home/cloudera/employee.txt /inputdirectroy
```

```
hdfs dfs -ls /
```

```
hdfs dfs -ls /inputdirectory
```

```
hadoop fs -cat /inputdirectory/employee.txt
```

```
hive
```

```
show databases;
```

```
create database organization;
```

```
show databases;
```

```
use organization;
```

```
show tables;
```

```
hive> create table employee(
```

```
> id int,
```

```
> name string,
```

```
> city string,
```

```
> department string,
```

```
> salary int,
```

```
> domain string)
```

```
> row format delimited
```

```
> fields terminated by '~';
```

```
show tables;
```

```
select * from employee;
```

```
show tables;
```

```
load data inpath '/inputdirectory/employee.txt' overwrite into table employee;
```

```
show tables;
```

```
select * from employee;
```

## PRACTICAL 6

### PRACTICAL 6A

#### Aim:

a) Implement Decision Tree classification technique using Social\_Network\_Ads.csv dataset.

#### Code:

```
# Decision Tree Classification
```

```
# Importing the dataset
```

```
dataset = read.csv("C:\\2022-23\\BDA PRACTICAL 2023\\Social_Network_Ads.csv")
```

```
#print(dataset)
```

```
dataset = dataset[3:5] # columns 3 4 ad 5
```

```
print(dataset)
```

```
# Encoding the target feature as factor(just like a vector having levels
```

```
# levels to convey that only two possible values for purchased - 0 & 1
```

```
dataset$Purchased = factor(dataset$Purchased, levels = c(0, 1))
```

```
print (dataset$Purchased)
```

```
# Splitting the dataset into the Training set and Test set
```

```
install.packages('caTools')
```

```
library(caTools)
```

```
set.seed(123)
```

```
#split = sample.split(dataset$Purchased, SplitRatio = 0.75)
```

```
split = sample.split(dataset$Purchased, SplitRatio = 0.75)
```

```
training_set = subset(dataset, split == TRUE)
```

```
test_set = subset(dataset, split == FALSE)
```

```
# Feature Scaling - scale() method centers and/or scales the columns of a numeric matrix.
```

```
training_set[-3] = scale(training_set[-3]) # scaling first 2 columns, don't consider 3rd column
```

```
test_set[-3] = scale(test_set[-3])
```

```
#print(test_set[-3])
```

```
# Fitting Decision Tree Classification to the Training set
```

```
install.packages('rpart')
```

```
library(rpart) # for partitioning tree
```

```
install.packages('rpart.plot')
```

```
library(rpart.plot)
```

```
classifier = rpart(formula = Purchased ~ ., data = training_set)
```

```
# Predicting the Test set results
```

```
y_pred = predict(classifier, newdata = test_set[-3], type = 'class')
```

```
print(y_pred)
```

```
# Making the Confusion Matrix
```

```
cm = table(test_set[, 3], y_pred)
print(cm)
```

```
y_grid = predict(classifier, newdata = grid_set, type = 'class')
```

```
# Plotting the tree
```

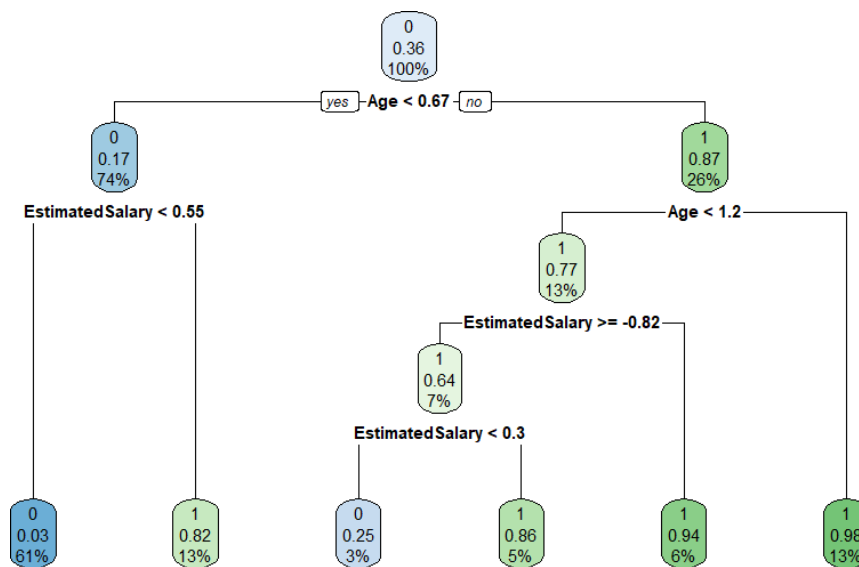
```
#extra=106 class model with a binary response
```

```
#extra=104 class model with a response having more than two levels
```

```
rpart.plot(classifier, extra = 106)
```

## Output:

```
> cm = table(test_set[, 3], y_pred)
> print(cm)
  y_pred
    0    1
0  53  11
1   6  30
```



## PRACTICAL 6B

### SVM Classification

#### Aim:

b) Implement SVM Classification technique using Social\_Network\_Ads.csv dataset. Evaluate the performance of classifier.

#### Code:

```
# Support Vector Machine (SVM)

# Importing the dataset
dataset = read.csv('C:\\2022-23\\BDA PRACTICAL 2023\\Social_Network_Ads.csv')
dataset = dataset[3:5]
print(dataset)
print(dataset$Purchased)
# Splitting the dataset into the Training set and Test set
install.packages('caTools')
library(caTools)
set.seed(123)
split = sample.split(dataset$Purchased, SplitRatio = 0.75)
training_set = subset(dataset, split == TRUE)
print(training_set)
test_set = subset(dataset, split == FALSE)
print(test_set)
# Feature Scaling
training_set[-3] = scale(training_set[-3]) # [-3] means 3rd index will be dropped
test_set[-3] = scale(test_set[-3])
print(training_set[-3])
print(test_set[-3])
# Fitting SVM to the Training set
install.packages('e1071')
library(e1071)
classifier = svm(formula = Purchased ~ .,
                 data = training_set,
                 type = 'C-classification',
                 kernel = 'linear')
print(classifier)

# Predicting the Test set results
y_pred = predict(classifier, newdata = test_set[-3])
print(y_pred)

# Making the Confusion Matrix
cm = table(test_set[, 3], y_pred)
print(cm)
```

#### Output:

```
> cm = table(test_set[, 3], y_pred)
> print (cm)
      y_pred
      0    1
0  57    7
1  13   23
```



## PRACTICAL 7

### PRACTICAL 7A

#### Logistic Regression

**Aim:**

a) Import data from web storage – binary.csv. Name the dataset and do Logistic Regression to find out relation between variables that are affecting the admission of a student in an institute based on his or her GRE score, GPA obtained and rank of the student. Also check the model is fit or not.

**Code:**

```
#fetch the data
college <- read.csv("https://raw.githubusercontent.com/csquared/udacity-
dlnd/master/nn/binary.csv")
head(college)
nrow(college)

install.packages("caTools") # For Logistic regression
library(caTools)

split <- sample.split(college, SplitRatio = 0.75)
split

training_reg <- subset(college, split == "TRUE")
test_reg <- subset(college, split == "FALSE")

# Training model
fit_logistic_model <- glm(admit ~ .,
                          data = training_reg,
                          family = "binomial")

# Predict test data based on model
predict_reg <- predict(fit_logistic_model,
                      test_reg, type = "response")
predict_reg

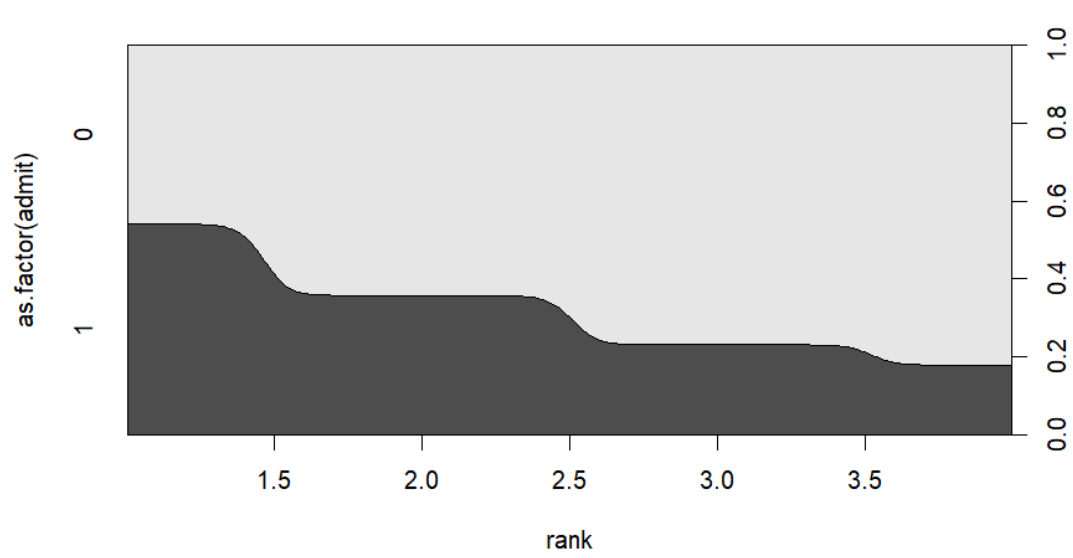
cdplot(as.factor(admit)~ gpa, data=college)
cdplot(as.factor(admit)~ gre, data=college)
cdplot(as.factor(admit)~ rank, data=college)

# Changing probabilities
predict_reg <- ifelse(predict_reg > 0.5, 1, 0)
predict_reg

# Evaluating model accuracy
# using confusion matrix
table(test_reg$admit, predict_reg)
```

### Output:

```
> table(test_reg$admit, predict_reg)
predict_reg
 0  1
0 70  2
1 21  7
```



## PRACTICAL 7B

### Multiple Regression Model

#### Aim:

b) Apply multiple regressions, if data have a continuous independent variable. Apply on above dataset – binary.csv.

#### Code:

```
#fetch the data
college <- read.csv("https://raw.githubusercontent.com/csquared/udacity-
dlnd/master/nn/binary.csv")
head(college)
nrow(college)

install.packages("caTools") # For Logistic regression
library(caTools)

split <- sample.split(college, SplitRatio = 0.75)
split


training_reg <- subset(college, split == "TRUE")
test_reg <- subset(college, split == "FALSE")

# Training model
fit_MRegressor_model <- lm(formula = admit ~ gre+gpa+rank,
                           data = training_reg)

# Predict test data based on model
predict_reg <- predict(fit_MRegressor_model,
                      newdata = test_reg)
predict_reg

cdplot(as.factor(admit)~ gpa, data=college)
cdplot(as.factor(admit)~ gre, data=college)
cdplot(as.factor(admit)~ rank, data=college)
```

#### Output:



## PRACTICAL 8

### PRACTICAL 8A

#### Classification Model Naïve Bayes Classification

**Aim:**

a) Implement Naïve Bayes Classification technique using Social\_Network\_Ads.csv dataset. Evaluate the performance of classifier.

**Code:**

```
# Naive Bayes
# Importing the dataset
dataset = read.csv('C:\\2022-23\\BDA PRACTICAL 2023\\Social_Network_Ads.csv')
dataset = dataset[3:5]
# Encoding the target feature as factor
dataset$Purchased = factor(dataset$Purchased, levels = c(0, 1))
# Splitting the dataset into the Training set and Test set
#install.packages('caTools')
library(caTools)
set.seed(123)
split = sample.split(dataset$Purchased, SplitRatio = 0.75)
training_set = subset(dataset, split == TRUE)
test_set = subset(dataset, split == FALSE)
# Feature Scaling
training_set[-3] = scale(training_set[-3])
test_set[-3] = scale(test_set[-3])
# Fitting Naive Bayes to the Training set
install.packages('e1071')
library(e1071)
classifier = naiveBayes(x = training_set[-3],
                        y = training_set$Purchased)
# Predicting the Test set results
y_pred = predict(classifier, newdata = test_set[-3])
# Making the Confusion Matrix
cm = table(test_set[, 3], y_pred)
print(cm)
```

**Output:**

```
> cm = table(test_set[, 3], y_pred)
> print(cm)
  y_pred
  0  1
0 57  7
1  7 29
```

## PRACTICAL 8B

### K-Means Clustering

#### Aim:

a) Clustering algorithms for unsupervised classification. Read a datafile Mall\_Customers.csv and apply k-means clustering. Plot the cluster data using R visualizations.

#### Code:

```
# K-Means Clustering

# Importing the dataset
dataset = read.csv('C:\\2022-23\\BDA PRACTICAL 2023\\Mall_Customers.csv')
head(dataset)
dataset = dataset[4:5]
head(dataset)


wcss = vector()
for (i in 1:10) wcss[i] = sum(kmeans(dataset, i)$withinss)
plot(1:10,
     wcss,
     type = 'b',
     main = paste('The Elbow Method'),
     xlab = 'Number of clusters',
     ylab = 'WSS')

# Fitting K-Means to the dataset with no of clusters = 5
kmeans = kmeans(x = dataset, centers = 5)
y_kmeans = kmeans$cluster

# Visualising the clusters
library(cluster)
clusplot(dataset,
         y_kmeans,
         lines = 0,
         shade = TRUE,
         color = TRUE,
         labels = 2,
         main = paste('Clusters of customers'),
         xlab = 'Annual Income',
         ylab = 'Spending Score')
```

#### Output:

### The Elbow Method



### Clusters of customers

